

Ralf Rantzau

[linkedin.com/in/rantzau](https://www.linkedin.com/in/rantzau)

Summary

Passionate about building systems in these areas:

- **Databases / SQL**
- **Data streams / Real-time analytics**
- **Cloud security**
- **Blockchain / DLT**
- **ML / Analytics / OLAP**
- **Distributed systems**

... and working with great people to build game-changing software platforms and solutions by

- Programming
- Inventing
- Research

Experience

Senior Software Engineer — Backend | Passio, Menlo Park | 12/2020-2/2021

- Enhanced and redesigned ingestion pipeline of multimedia data for Edge AI solutions
- Enhanced backend APIs

Senior Software Engineer — Data platform | Lacework, San Jose | 6/2020-11/2020

- Enhanced query processing feature for the cloud security data pipeline

Tech Lead — Blockchain platform | Cisco, San Jose | 2018-9/2019

- [Cisco Blockchain](#) — Platform backend programming
- Designed, built, and delivered industry-first DLT platform features, e.g., services to enable role-based information disclosure (field-level RBAC security)
- Worked on the implementation of the state store and ledger
- Patents

Tech Lead — Cloud telemetry and log analytics | Cisco, San Jose | 2014-2018

- [Cisco Zeus](#) — Database backend programming, built BI reports and tools, telemetry & logs crunching with Kafka, Storm, Cassandra
- [Cisco Container Platform](#) — Built infrastructure to capture telemetry & logs
- [Cisco Cloud Center](#) — Built infrastructure to capture telemetry and logs using containers
- Delivered analytics micro-services for product usage optimization and management
- Delivered product features based on stochastic streaming algorithms over telemetry data (count-min sketch, hyperloglog sketch)
- Managed team to deploy a multi-cloud version of Zeus with Spinnaker as well as Cisco Cloud Center
- Managed team to build a telemetry router for logs, metrics, alerts, and health checks
- Designed and built backend infrastructure for processing and metrics/time-series data

- Patents

Senior Software Engineer — Real-time analytics | Cisco, San Jose | 2012-2014

- Cisco Connected Streaming Analytics (CSA) — Stream engine programming in C
- Customer PoCs with Cisco CSA
- ML algorithms integration into real-time analytics pipeline
- Edge and fog computing — Analytics for IoT (retail, stadiums)
- Delivered customer solutions based on Connected Streaming Analytics based with machine learning
- Predicted and optimized retail shopper checkout queues
- Optimized retail store layout based on purchase patterns and shopper tracking
- Integrated machine learning algorithms into stream database queries using R and user-defined-functions
- Patents & publications

Assistant Professor — Database systems | Aalborg University, Denmark | 2010-2011

- Internet technologies lecture
- Database introduction lecture
- Renewable energy data management research

Software Engineer — Data streams & real-time query processing | Truviso, Foster City | 2009-2012

- [Truviso TruCQ](#) — Stream engine programming in C (think of PostgreSQL with data streams)
- PostgreSQL query optimization
- Built PostgreSQL database kernel features for real-time stream processing such as sessionization
- **Acquired by Cisco** in 2012

Senior Software Engineer — RFID data management & BI | IBM Silicon Valley Lab, San Jose | 2006-2009

- IBM RFID Information Center — Database backend and middleware programming in Java
- IBM Cognos — Data governance in data warehouses, SQL query optimization
- Patents & publications

Post-doctoral Scientist — Data privacy technology | IBM Almaden Research Center, San Jose | 2004-2006

- Hippocratic Database — Enforcing privacy policies in relational database (DBMS agnostic)
- Graph databases — Privacy technology for large multi-media streams
- Search engines — Context-sensitive ranking of web search results
- Patents & publications

Teaching Assistant — DBMS implementation, BI, Data mining | University of Stuttgart, Germany | 1999-2003

- Data mining / analytics — Algorithms, integration into database engine for high-performance crunching
- Data warehousing / business intelligence (BI) — Database design, query optimization, performance tuning
- Databases — Architecture, implementation, SQL, object-relational technology
- IOT / Smart Factories — IT for digital production processes

Assistant Lecturer | Stuttgart Institute of Management and Technology, Germany | 1998-2000

- Database systems lectures

Education

Ph.D. Computer Science | University of Stuttgart, Germany | 2001-2004

- Dissertation: [Query Processing Concepts and Techniques for Set Containment Tests](#)
- [Database Systems Department](#) (IPVS/AS), Prof. Dr.-Ing. habil. Bernhard Mitschang

M.S. Computer Science | University of Stuttgart, Germany | 1991-1997

- Master thesis: [Extended Concepts for Association Rule Discovery](#)
- [Database Systems Department](#) (IPVS/AS), Prof. Dr. rer. nat. Dr. h. c. Kurt Rothermel

Bachelor Business Administration | University of Stuttgart, Germany | 1993-1996

- Strategic Management
- Production Planning
- System Dynamics
- Accounting
- Business IT
- People Management

Certifications

- [Blockchain Basics](#) by University at Buffalo & The State University of New York, Coursera, 2/2020
- [Neural Networks and Deep Learning](#), Coursera, 1/2018
- [Improving Deep Neural Networks: Hyperparameter tuning, Regularization and Optimization](#), Coursera, 1/2018
- [Serverless Data Analysis with Google BigQuery and Cloud Dataflow](#), Coursera, 1/2018

Awards

- Best Paper Award: [Laws for Rewriting Queries Containing Division Operators](#) (Ralf Rantza, Christoph Mangold), [22nd International Conference on Data Engineering \(ICDE\)](#), Atlanta, Georgia, USA, 2006
- Patent Achievement Award, IBM, 2009

Skills

According to LindedIn endorsements ...

- Big Data
- Databases
- Distributed Systems
- Software Development
- Software Engineering

Languages

- C/C++
- Golang (preferred)
- Java
- Python
- SQL
- Experience with R

Patents & Disclosures

1. [Auditing compliance with a Hippocratic database](#)
2. [Security and privacy enforcement for discovery services in a network of electronic product code information repositories](#)
3. [Access control method and a system for privacy protection](#)
4. [Mitigating and managing privacy risks using planning](#)

5. Providing security for queries to electronic product code information services
6. Entity-state relationship correlation
7. Multi-datacenter message queue
8. Efficient trickle updates in large databases using persistent memory
9. Inter-tenant workload performance correlation and recommendation
10. Using persistent memory to enable restartability of bulk load transactions in cloud databases
11. Traffic analytics service for telemetry routers and monitoring systems
12. Using persistent memory to enable consistent data for batch processing and streaming processing
13. Chained collection of information
14. Obfuscation and anonymization techniques for network data sets for machine learning
15. Enforcing network endpoint policies in a cloud-based environment using a covert namespace
16. Self-organizing meshes using distributed ledger systems
17. Hop-by-hop authentication routing using distributed ledger systems
18. Dynamic Genesis and cold death of distributed ledger systems
19. Intent-based ad-hoc network meshes using distributed ledger systems
20. Group addressing using distributed ledger systems
21. Onion routing for obfuscating traffic using distributed ledger systems
22. Ad-hoc n-tier global overlay using distributed ledger systems
23. Public key based addressing using distributed ledger systems
24. Supply chain device provenance

Publications

1. Enhancing the functionality of the web, Peter Kutschera and Ralf Rantza. ADBIS, September 1997.
2. Extended concepts for association rule discovery. Ralf Rantza. Master thesis no. DIP-1554, University of Stuttgart, Faculty of Computer Science, Electrical Engineering, and Information Technology, December 1997.
3. A multi-tier architecture for high-performance data mining. Ralf Rantza and Holger Schwarz. In Proceedings BTW, Freiburg, Germany, March 1999.
4. StreamJoin: A generic database approach to support the class of stream-oriented applications. Clara Nippl, Ralf Rantza, and Bernhard Mitschang. In Proceedings IDEAS, Yokohama, Japan, September 2000.
5. SIES — An approach for a federated information system in manufacturing. Carmen Constantinescu, Uwe Heinkel, Ralf Rantza, and Bernhard Mitschang. In Proceedings International Symposium on Information Systems and Engineering (ISE), Las Vegas, Nevada, USA, June 2001.
6. Universal quantification in relational databases: A classification of data and algorithms. Ralf Rantza, Leonard Shapiro, Bernhard Mitschang, and Quan Wang. In Proceedings EDBT, Prague, Czech Republic, March 2002.
7. Frequent itemset discovery with SQL using universal quantification. Ralf Rantza. In Proceedings DTD, Prague, Czech Republic, March 2002.
8. A system for data change propagation in heterogeneous information systems. Carmen Constantinescu, Uwe Heinkel, Ralf Rantza, and Bernhard Mitschang. In Proceedings ICEIS, Ciudad Real, Spain, April 2002.
9. Champagne: Data change propagation for heterogeneous information systems. Ralf Rantza, Carmen Constantinescu, Uwe Heinkel, and Holger Meinecke. In Proceedings VLDB, software demonstration, Hong Kong, China, August 2002.
10. Query processing concepts and techniques to support business intelligence applications. Ralf Rantza. In Proceedings VLDB, doctoral poster session, Hong Kong, China, August 2002.
11. Algorithms and applications for universal quantification in relational databases. Ralf Rantza, Leonard Shapiro, Bernhard Mitschang, and Quan Wang. Information Systems, 28(1), January 2003.
12. Föderal: Management of engineering data using a semistructured data model. Christoph Mangold, Ralf Rantza, and Bernhard Mitschang. In Proceedings ICEIS, Angers, France, April 2003.

13. Processing frequent itemset discovery queries by division and set containment join operators. Ralf Rantza. In Proceedings DMKD, San Diego, California, USA, June 2003.
14. Divide et impera: A flexible integration of layout planning and logistics simulation through data change propagation. Bernhard Mitschang, Engelbert Westkämper, Carmen Constantinescu, Uwe Heinkel, Benno Löffler, Ralf Rantza, and Ralph Winkler. In Proceedings CIRP International Seminar on Manufacturing Systems, Saarbrücken, Germany, June 2003.
15. Optimierung von Abfragesequenzen in Business-Intelligence-Anwendungen. Holger Schwarz, Tobias Kraft, Ralf Rantza, and Bernhard Mitschang. Information Technology, 2003(4), August 2003.
16. Coarse-grained optimization: Techniques for rewriting SQL statement sequences. Tobias Kraft, Holger Schwarz, Ralf Rantza, and Bernhard Mitschang. In Proceedings VLDB, Berlin, Germany, September 2003.
17. Frequent itemset discovery with SQL using universal quantification. Ralf Rantza. In Rosa Meo, Pier L. Lanzi, and Mika Klemettinen, editors, Database support for data mining applications, volume 2682 of LNCS. Springer, 2004. ISBN: 3-540-22479-3.
18. Query processing concepts and techniques for set containment tests. Ralf Rantza. Ph.D. thesis no. 2004/01, University of Stuttgart, Faculty of Computer Science, Electrical Engineering, and Information Technology, January 2004. Also published as book, ISBN 978-3-639-24439-7, VDM Verlag.
19. Auditing compliance with a Hippocratic database. Rakesh Agrawal, Roberto J. Bayardo Jr., Christos Faloutsos, Jerry Kiernan, Ralf Rantza, and Ramakrishnan Srikant. In Proceedings VLDB, Toronto, Canada, August-September 2004.
20. Laws for rewriting queries containing division operators. Ralf Rantza and Christoph Mangold. Technical report no. 2005/08, University of Stuttgart, Faculty of Computer Science, Electrical Engineering, and Information Technology, October 2005.
21. Laws for rewriting queries containing division operators. Ralf Rantza and Christoph Mangold. In Proceedings ICDE, Atlanta, Georgia, USA, April 2006. Best paper award.
22. Context-sensitive ranking. Rakesh Agrawal, Ralf Rantza, and Evimaria Terzi. In Proceedings SIGMOD, Chicago, Illinois, USA, June 2006.
23. Discovery services—Enabling RFID traceability in EPCglobal networks (demonstration paper). Steve Beier, Tyrone Grandison, Karin Kailing, and Ralf Rantza. In Proceedings COMAD, Delhi, India, December 2006.
24. Theseos query engine for traceability networks. Tyrone Grandison, Karin Kailing, and Ralf Rantza. In IBM alphaWorks, January 2007.
25. Security with WebSphere RFID Information Center 1.0. Ralf Rantza and Dirk Wollscheid. In IBM developerWorks, March 2007.
26. An alert notification facility for RFID event data repositories. Valer-Alin Crisan and Ralf Rantza. In Proceedings RFID Data Management (RFDM), Cancun, Mexico, April 2008.
27. Cost-based predictive spatio-temporal join. Wook-Shin Han, Jaehwa Kim, Byung Suk Lee, Yufei Tao, Ralf Rantza, and Volker Markl. Transactions on Knowledge and Data Engineering, 28(2), February 2009.